Gagan Sharma

gagansharma3002@gmail.com | +91 96443 65693 | LinkedIn | GitHub | LeetCode

SUMMARY

Software engineer specializing in AI infrastructure, distributed systems, and production AI platforms. Focused on architecting systems where retrieval, execution, and orchestration decisions directly reduce latency, cut inference cost, and improve reliability at scale — including a 60% reduction in AI inference cost and compressing multi-hour evaluation cycles to under 3 minutes through distributed execution.

EXPERIENCE

Software Engineer (AI Platform) Oct 2025 – Jul 2026 *Appointy — Onsite*

GAP — Generative AI Platform

- Every new AI feature risked re-solving the same problems — retrieval, context management, tool access, and evaluation — from scratch, with no standard way to catch quality regressions before production. Architected a modular platform separating agent orchestration, RAG, memory, tool integrations, and evaluation into reusable components, letting new applications — including GCP Resource Efficiency Agent — compose existing infrastructure and validate output against defined criteria instead of rebuilding and re-testing from zero.

GCP Resource Efficiency Agent — Built on GAP

- Cloud teams lacked a reliable way to identify unused, underutilized, and overutilized GCP resources across projects, so operators spent hours cross-referencing dashboards and bills without a clear view of what to fix or why. Built a reactive intelligence layer, on top of GAP, that retrieves infrastructure history and answers natural-language questions — why costs rose, which resources are underutilized, how this week compares to last — with grounded explanations instead of raw charts.
- Cost and capacity issues were typically discovered only after waste had already accumulated, and recommendations were easy to ignore if engineers had to leave their workflow to investigate them. Designed a proactive layer that continuously scans configured GCP projects, ranks optimization opportunities by estimated savings, effort, and risk, and surfaces them directly in Slack with supporting evidence — shifting infrastructure management from reactive investigation to prioritized, actionable weekly reports.

Backend Developer Intern Feb 2025 – Jul 2025 *Digital Live 24 — Remote*

- Growing booking and content-management volume on the FlyShuttle and Digital Signage platforms exposed the limits of ad hoc backend logic under production load. Consolidated core workflows onto a backend architecture built for transactional integrity, allowing both platforms to scale with customer demand without service disruption.
- Notification and data-sync operations were tightly coupled to the main request cycle, so a delay in one workflow risked cascading failures across unrelated transactions. Decoupled these operations through an event-driven architecture to isolate failures, improving platform reliability under concurrent load.
- Inconsistent access control and slow endpoints limited how much traffic the platforms could serve reliably, capping growth. Centralized role-based access control and optimized the slowest API paths, improving both security posture and response times as customer usage scaled.

Backend Developer Intern Sep 2024 – Feb 2025 *Katyayani Organics — Hybrid*

- Procurement, inventory, QC, and production tracking ran across disconnected manual processes, slowing cross-team visibility and decision-making. Unified these workflows onto a single backend platform with a shared data model, improving API performance by 30% and giving teams one consistent source of truth.

PROJECTS

Pnyx — Real-Time Meeting Intelligence Platform

- Retrieving complete meeting history for every AI query increased latency and inference cost as conversations accumulated. Architected a hierarchical retrieval strategy scoping context to meeting, workspace, and cross-meeting layers instead of searching full history, balancing response quality with token efficiency — reducing token consumption by 60%.
- Synchronous processing of meeting audio and AI responses risked blocking user-facing requests as concurrent meeting volume grew. Introduced asynchronous execution using Celery and Redis, decoupling long-running AI workloads from the request path and streaming transcription over WebSockets — delivering sub-2-second latency without sacrificing concurrency.

Ohara Benchmark — Distributed LLM Evaluation Framework

- Sequential evaluation of coding agents against private repositories took hours per cycle, slowing iteration on agent improvements. Designed a distributed, parallelized benchmark architecture that traded single-node simplicity for horizontal throughput, cutting evaluation time from hours to under 3 minutes.
- Manual grading was inconsistent, and public benchmark data risked contamination from prior model training. Built a blind LLM judging system with contamination-safe evaluation, sourcing real task context from GitHub, GitLab, and Jira instead of static datasets — producing repeatable, defensible scores against live engineering workflows.

TECHNICAL SKILLS

Languages: Python, Go, JavaScript, TypeScript, Java
AI Engineering: RAG, AI Agents, Agent Orchestration, Vector Search, LLM APIs, Prompt Engineering
Backend: FastAPI, Django, Node.js, Express, REST APIs, WebSockets, Microservices
Distributed Systems: Event-Driven Architecture, Celery, Kafka, Concurrency, Load Balancing, System Design
Databases: PostgreSQL, MongoDB, MySQL, Redis
Cloud & DevOps: Docker, Kubernetes, AWS, GCP, GitHub Actions, Linux, Nginx

EDUCATION

B.Tech in Computer Science 2026 *UIT-RGPV, Bhopal*